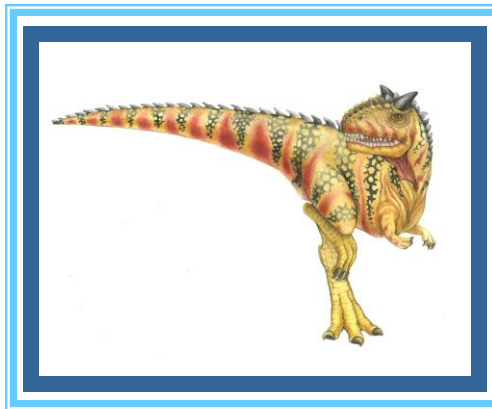


# Memory Management (contd.)

---





- **Slides mostly borrowed from Silberschatz & Galvin**
  - With occasional modifications from us





# The story so far

---

- **Basic motivation and concepts of memory management**
  - Address binding: process of mapping address space in the code to address space in main memory
  - Concept of logical/Virtual and physical address space
  - MMU: H/W that maps virtual to physical address
  - Swapping (context switch, double buffering)





# The story so far


---

- **Basic motivation and concepts of memory management**
  - Address binding: process of mapping address space in the code to address space in main memory
  - Concept of logical/Virtual and physical address space
  - MMU: H/W that maps virtual to physical address
  - Swapping (context switch, double buffering)
  
- **Allocating memory to processes**
  - Contiguous, multiple-partition
  - Problems: dynamic storage allocation, internal/external fragmentation
  - Segmentation: logical parts of a program, can be distributed throughout physical memory
  - Paging: Pages, frames, page table



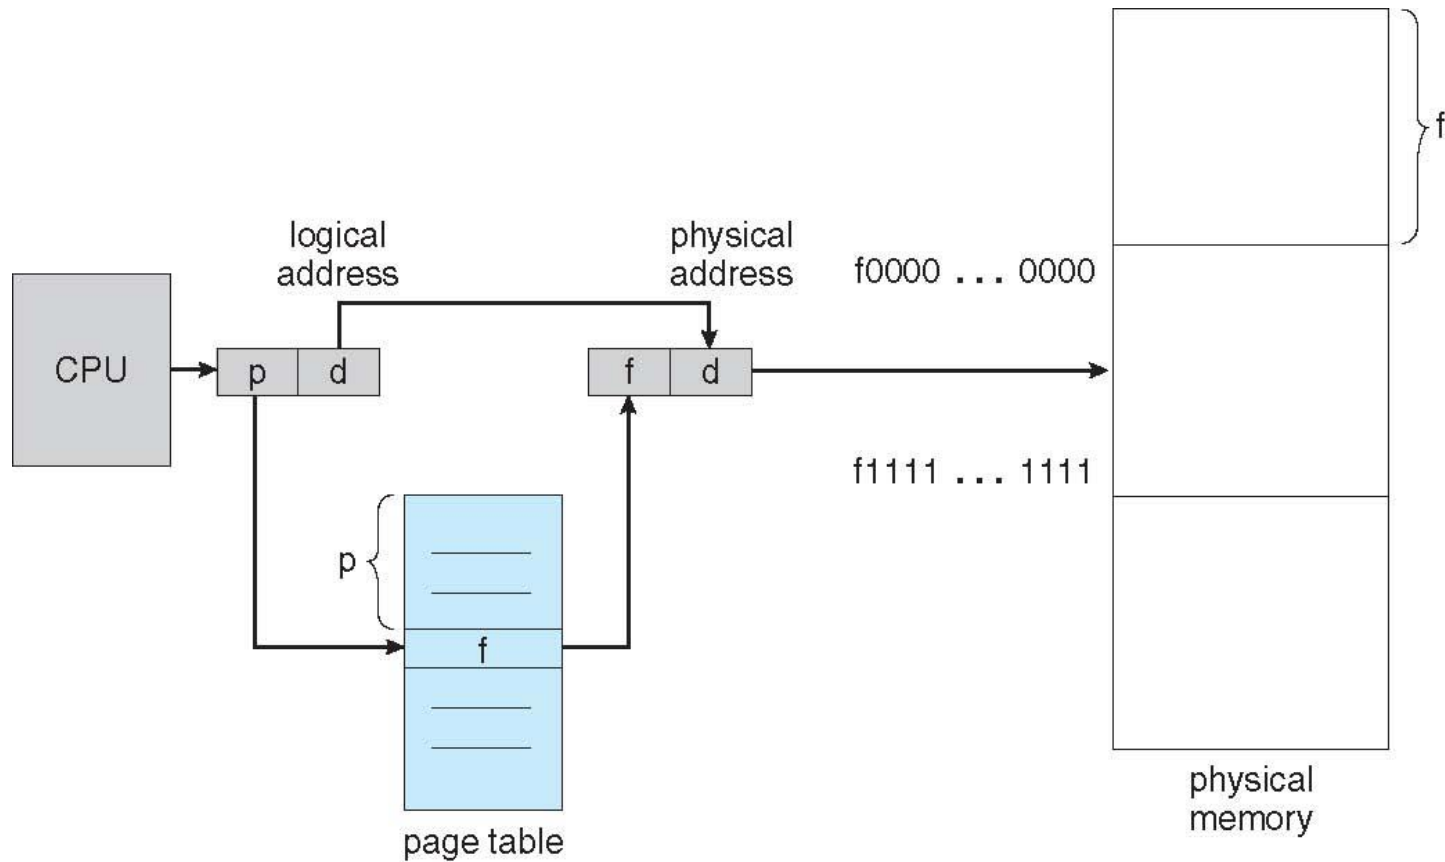


# Paging Model - Recap





# Paging Hardware - Recap





# Advantage of paging: Shared Pages

---

## □ Shared code

- One copy of read-only (**reentrant**) code shared among multiple processes (i.e., text editors, compilers, window systems)
- Also useful for inter-process communication if sharing of read-write pages is allowed

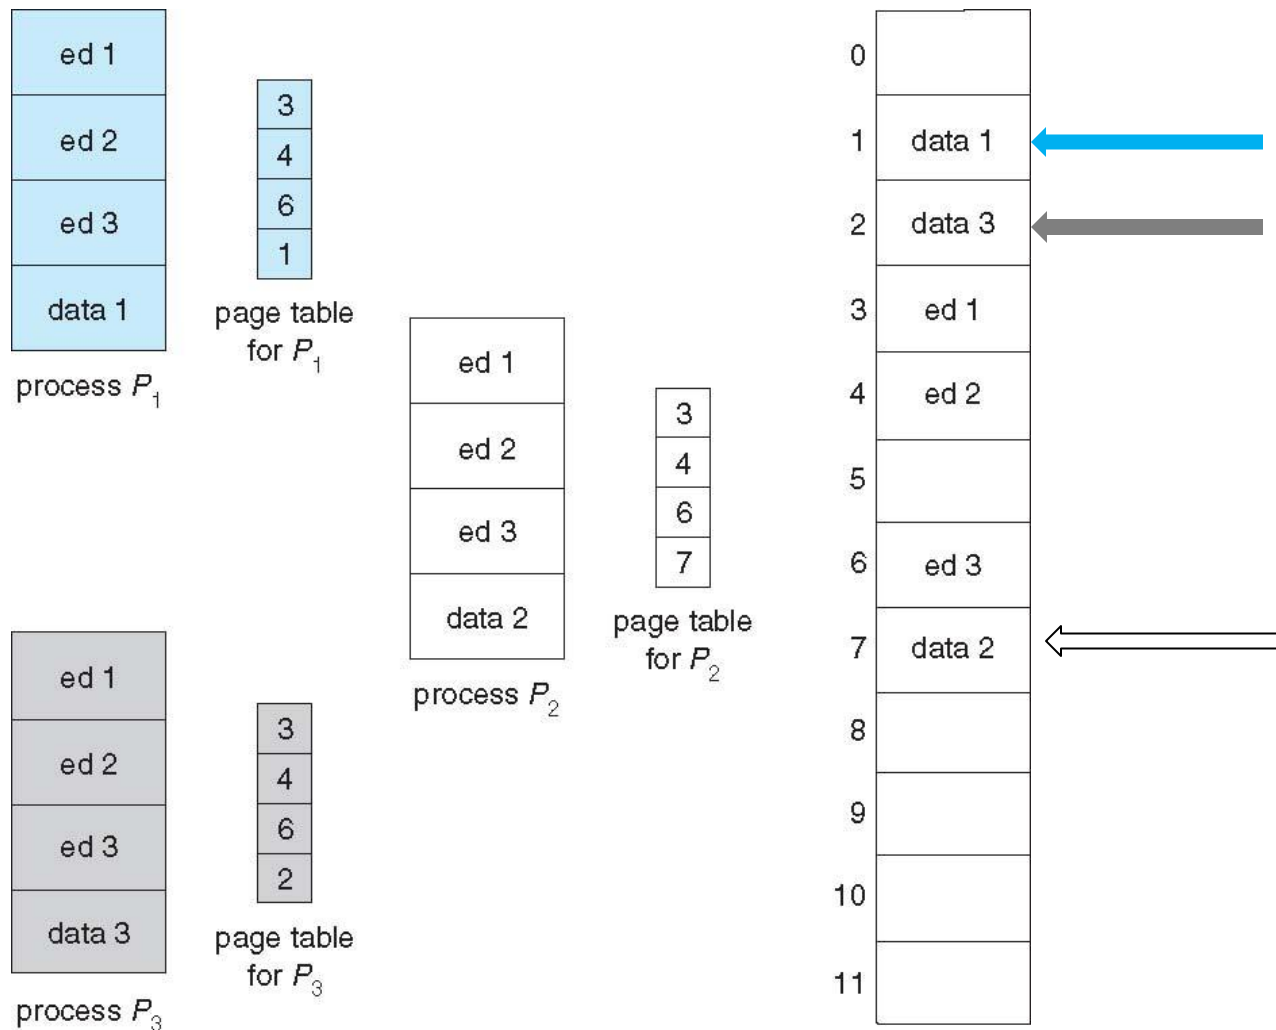
## □ Private code and data

- Each process might also keep its own copy of private code / data
- The pages for the private code / data can appear anywhere in the logical address space





# Shared Pages Example







# Memory Protection in Paging

- ❑ Memory protection implemented by associating **protection bits** with each page (these bits usually kept in the page table)
  - ❑ Indicates if read-only or read-write access is allowed
  - ❑ Can also add more bits to indicate page execute-only, and so on
- ❑ **Valid-invalid** bit attached to each entry in the page table:
  - ❑ “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal (valid) page
  - ❑ “invalid” indicates that the page is not in the process’ logical address space
  - ❑ Can use a **page-table length register (PTLR)** to indicate size of the page table, i.e., which pages are valid
- ❑ Any violations result in a trap to the kernel





# Valid (v) or Invalid (i) Bit In A Page Table

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	

frame number		valid-invalid bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page <i>n</i>





# Memory Management: Topics

---

- ❑ Background
- ❑ Swapping
- ❑ Contiguous Memory Allocation
- ❑ Segmentation
- ❑ Paging
- ❑ Structure of the Page Table (and variants of the page table)





# Implementation of Page Table

---

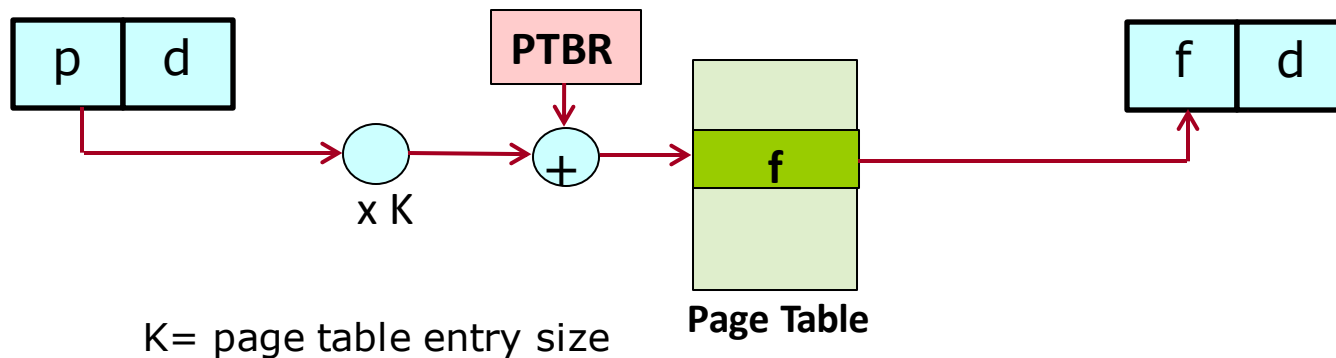
- Page Table has to be accessed for every memory access by a process
  - Implementation should allow fast access
- Where to store the Page Table?
- Best case – a set of dedicated registers
  - Only feasible if the page table is small
  - Not practical in most modern systems
- Practical case – **Page Table stored in the memory**





# Implementation of Page Table

- In practice, page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- In this scheme **every data/instruction access requires two memory accesses**
  - One for the page table entry, and one for the actual data/instruction
  - Too expensive





# Implementation of Page Table (Cont.)

- The two memory access problem can be solved by the use of a special fast-lookup **hardware cache** called **associative memory** or **translation look-aside buffers (TLBs)**
- TLBs typically small (between 32 and 1,024 entries)
  - Contains a few of the (most frequently/recently accessed) page table entries
  - For a logical address generated by CPU, its page number is searched in TLB; If page number is found (**TLB hit**), its frame number is immediately available to access the data directly in memory
- On a **TLB miss**, a memory reference to the page table needed
  - Value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access





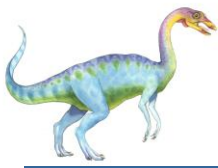
# Associative Memory

- TLB is associative memory – allows parallel search (item to be searched is compared with all keys simultaneously)


Page #	Frame #

- Address translation (p, d)
  - If p is in associative register, get Frame# out
  - Otherwise get Frame# from page table in memory





# Paging Hardware With TLB







# Effective Access Time

- Associative Lookup =  $\varepsilon$  time unit
  - Can be  $< 10\%$  of memory access time
- Hit ratio =  $\alpha$  (% of times that a page number is found in TLB)

- **Effective Access Time (EAT)**

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$

Even larger for multi-level page tables

- Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio  $\alpha = 99\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$





# Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers:  **$2^{32}$  byte memory locations**
  - Page size of 4 KB ( $4 \times 2^{10}$  byte =  $2^{12}$  byte)
  - **Page table** would have  **$\sim 1$  million entries** ( $2^{32} / 2^{12}$ )
  - If each page table entry is 4 bytes
    - ▶ 4 MB ( $2^{20} \times 4$  byte) of physical address space / memory for page table alone – equivalent to 1024 pages of 4KB each
    - ▶ That amount of memory used to cost a lot
    - ▶ **Don't want to allocate that contiguously in main memory**





# Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers:  **$2^{32}$  byte memory locations**
  - Page size of 4 KB ( $4 \times 2^{10}$  byte =  $2^{12}$  byte)
  - **Page table** would have  **$\sim 1$  million entries** ( $2^{32} / 2^{12}$ )
  - If each page table entry is 4 bytes
    - ▶ 4 MB ( $2^{20} \times 4$  byte) of physical address space / memory for page table alone – equivalent to 1024 pages of 4KB each
    - ▶ That amount of memory used to cost a lot
    - ▶ **Don't want to allocate that contiguously in main memory**

We moved to paging because we wanted to avoid allocating large blocks of contiguous memory

But we are back to the same problem – as logical address space increases, page tables need large memory blocks





# Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers:  **$2^{32}$  byte memory locations**
  - Page size of 4 KB ( $4 \times 2^{10}$  byte =  $2^{12}$  byte)
  - **Page table** would have  **$\sim 1$  million entries** ( $2^{32} / 2^{12}$ )
  - If each page table entry is 4 bytes
    - ▶ 4 MB ( $2^{20} \times 4$  byte) of physical address space / memory for page table alone – equivalent to 1024 pages of 4KB each
    - ▶ That amount of memory used to cost a lot
    - ▶ **Don't want to allocate that contiguously in main memory**
- **Solutions**
  - Hierarchical Paging
  - Hashed Page Tables
  - Inverted Page Tables





# Hierarchical Page Tables

---

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then **page the page table**






# Hierarchical Page Tables


---

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then **page the page table (POPT: pages of page table)**





# Two-Level Page-Table Scheme





# Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:

page number		page offset
$p_1$	$p_2$	$d$
12	10	10


- where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the inner page table
- Known as **forward-mapped page table**








# Address-Translation Scheme





# Address-Translation Scheme



Cost: Three memory  
Accesses  
(previously two)



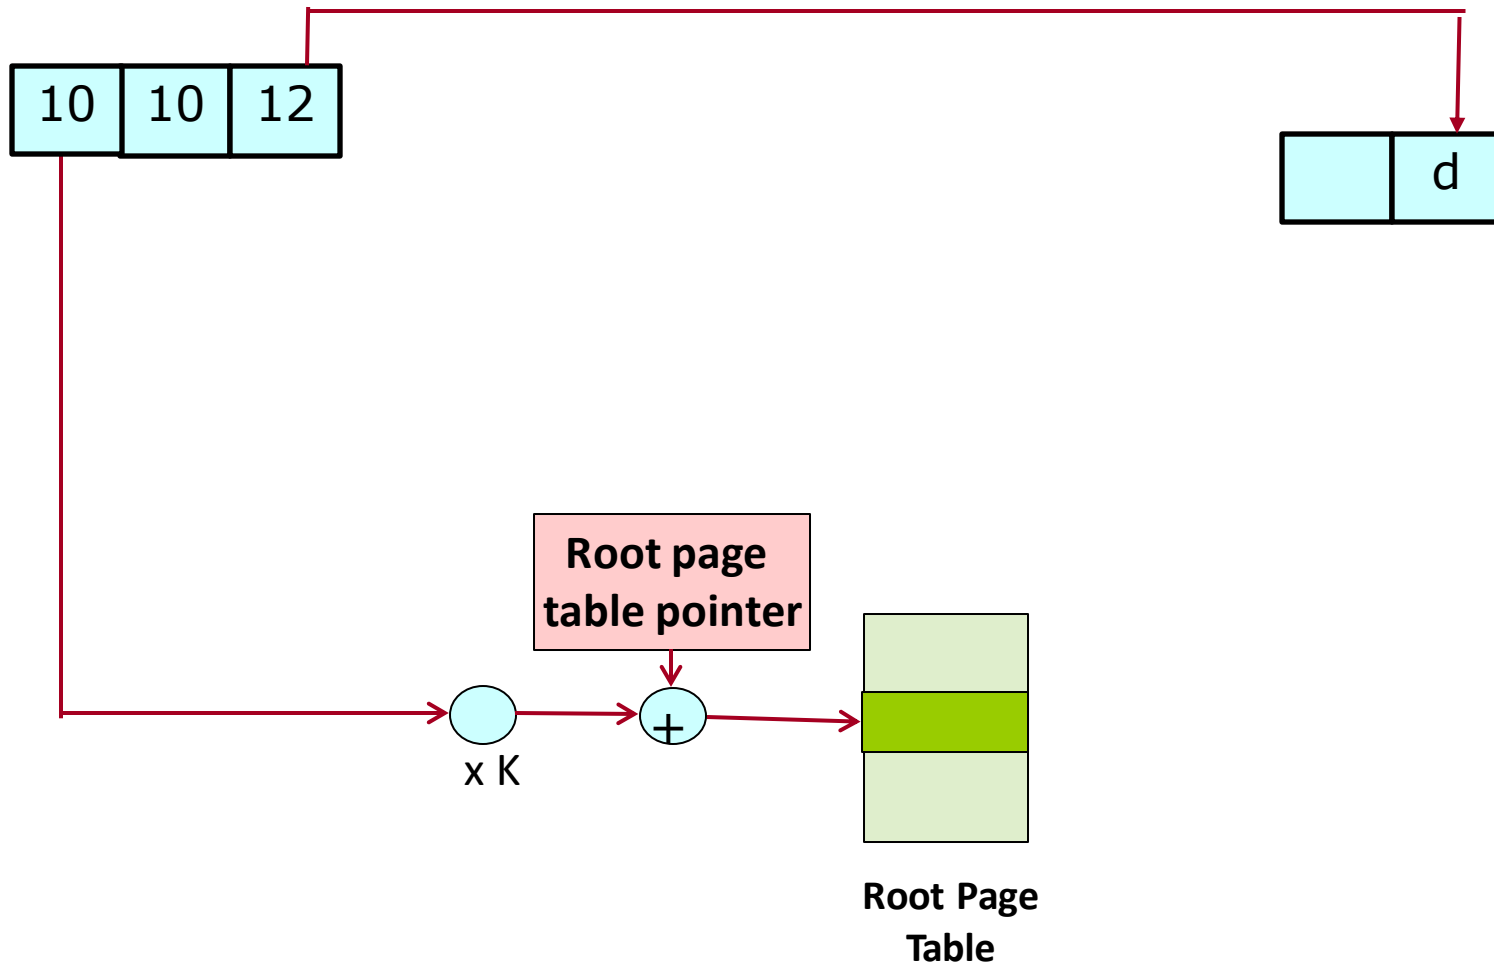


# Implementation of 2-level paging



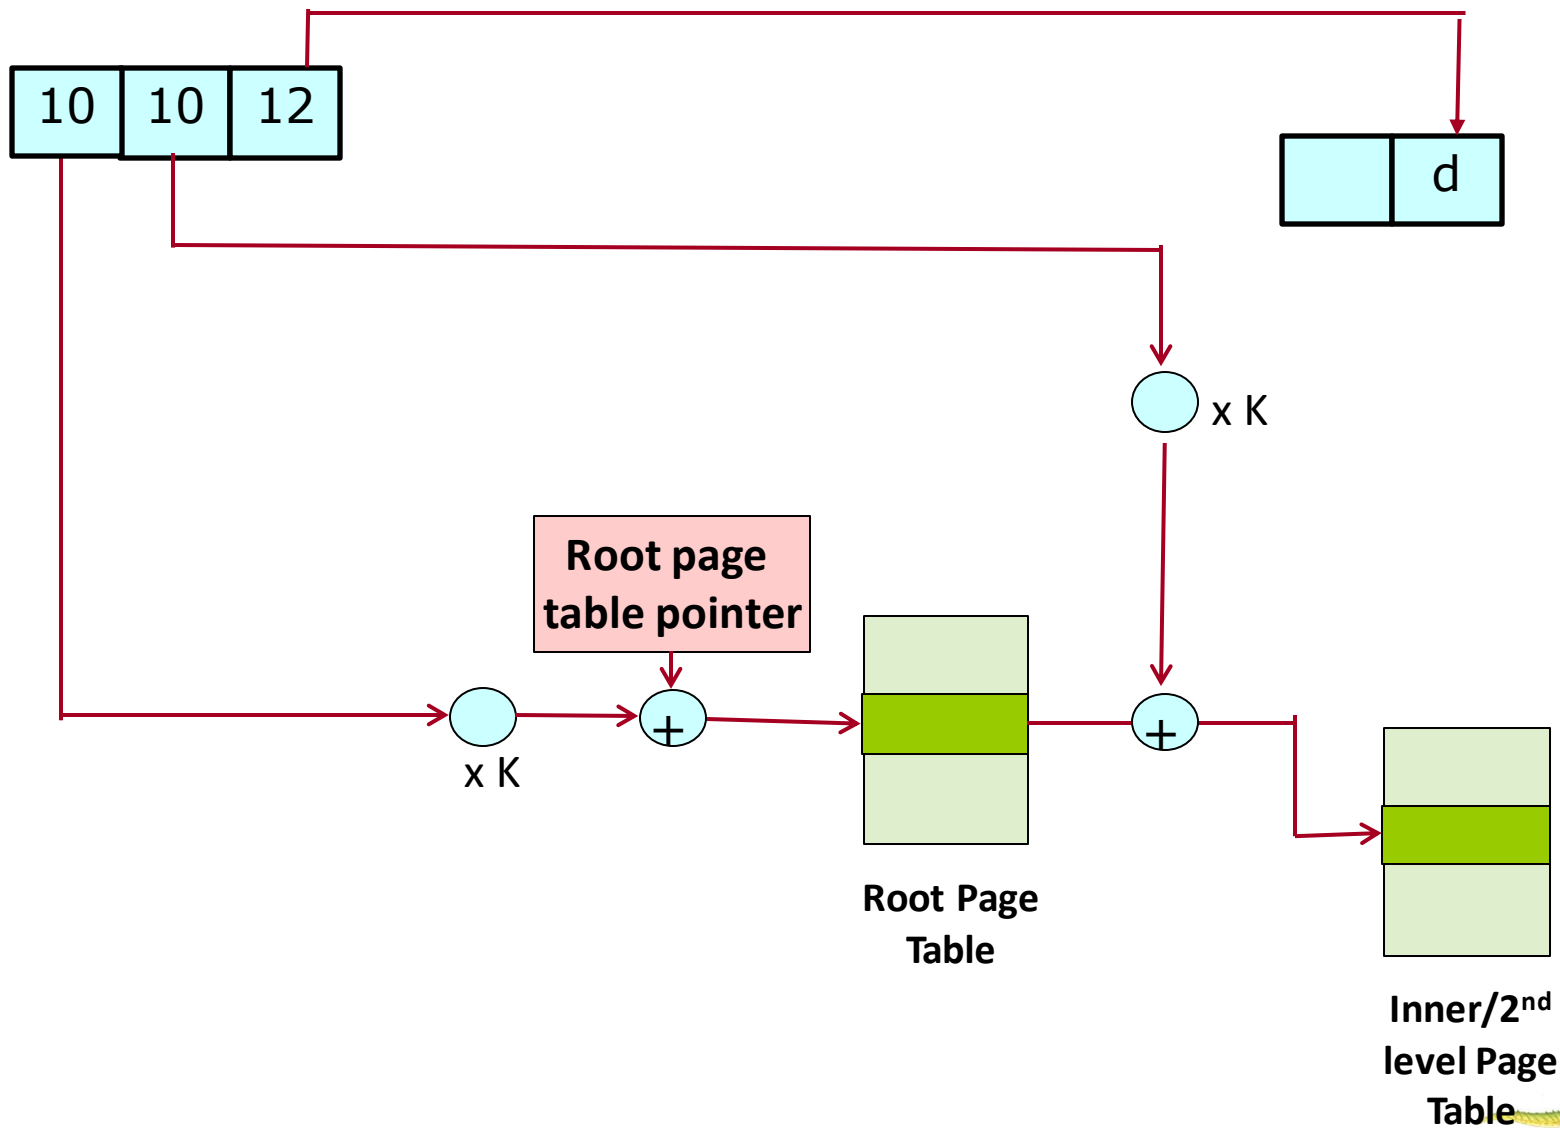


# Implementation of 2-level paging



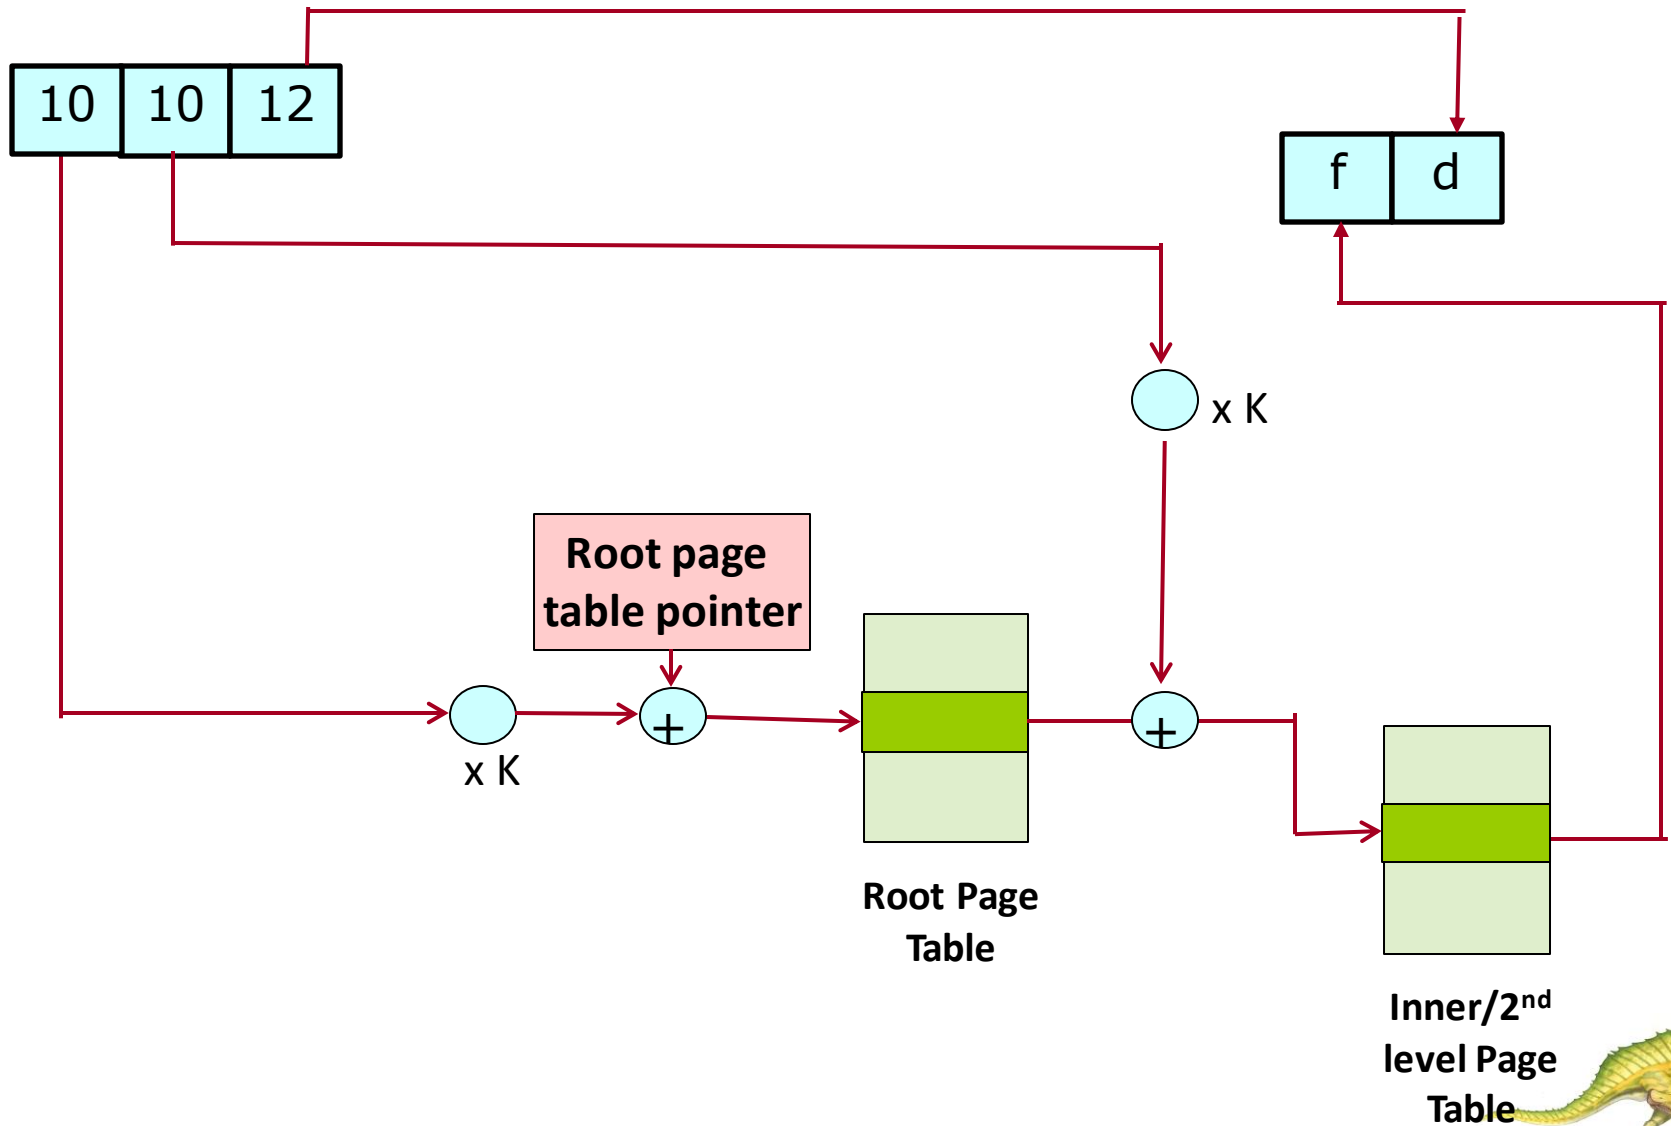


# Implementation of 2-level paging






# Implementation of 2-level paging





# 64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If 2-level scheme, inner/2<sup>nd</sup> level page tables could be  $2^{10}$  4-byte entries
  - Address would look like



- Outer page table has  $2^{42}$  entries or  $2^{44}$  bytes =  **$2^{14}$  GB**

Won't work!





# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

- One solution is to add a 2<sup>nd</sup> outer page table







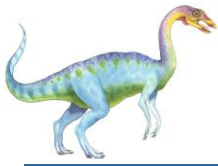
# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

- One solution is to add a 2<sup>nd</sup> outer page table

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12





# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

- One solution is to add a 2<sup>nd</sup> outer page table

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12

- But in this example the 2<sup>nd</sup> outer page table is still 2<sup>34</sup> bytes in size
  - ▶ Even 16 GB is too huge to store contiguously
  - ▶ Also 4 memory access to get to one physical memory location





# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

- One solution is to add a 2<sup>nd</sup> outer page table

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12

- But in this example the 2<sup>nd</sup> outer page table is still 2<sup>34</sup> bytes in size
  - ▶ Even 16 GB is too huge to store contiguously
  - ▶ Also 4 memory access to get to one physical memory location

Multi-level paging schemes do not scale well  
Need alternate mechanisms





# Hashed Page Tables

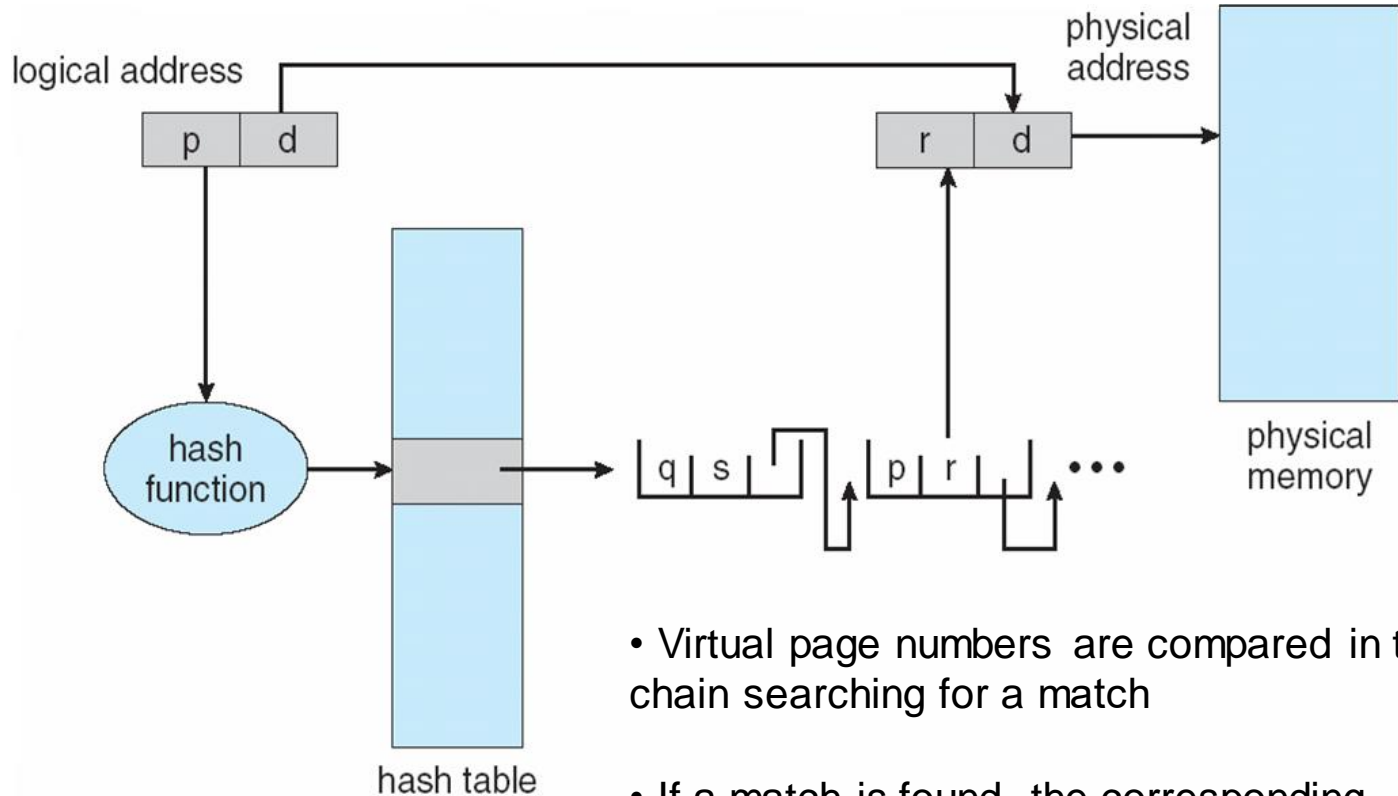
---

- Common in address spaces  $> 32$  bits
- The **virtual page number is hashed into a page table**
  - This page table contains a chain of elements hashing to the same location
- Each element contains
  - (1) the virtual page number
  - (2) the value of the mapped physical page frame
  - (3) a pointer to the next element





# Hashed Page Table



- Virtual page numbers are compared in this chain searching for a match
- If a match is found, the corresponding physical frame is extracted





# Inverted Page Table: Background

- Till now what we have discussed:
  - Each process has a page table that keeps track of all logical pages of this process

Process 1	X	Process 2	F4
	F1		X
	F3		X
	X		X
	X		F7
	F8		F11
	X		X



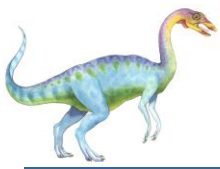


# Inverted Page Table

---

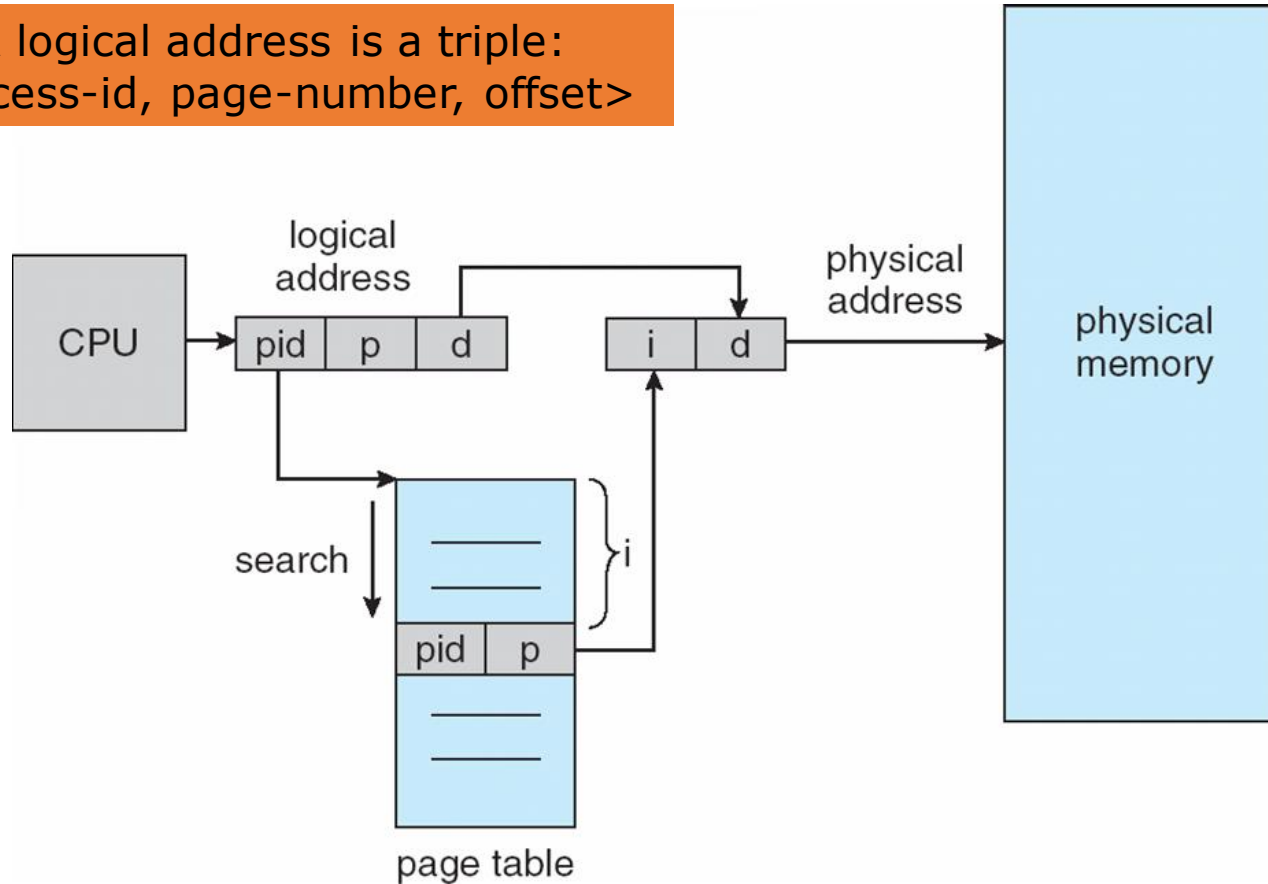
- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages (frames)
- **Inverted page table: One entry for each real page (frame) of memory**
  - Entry = the virtual address of the page stored in that real memory location + information of the process that owns the page





# Inverted Page Table Architecture

A logical address is a triple:  
<process-id, page-number, offset>







# Inverted Page Table

---

- ❑ Decreases memory needed to store each page table, but **increases time needed to search the table for a page reference**
  - ❑ Searching uses the CPU-generated virtual address, but the inverted page table is sorted by the physical address
- ❑ Can make searching physical address relatively fast:
  - ❑ Use hash table (as described earlier)
  - ❑ TLB can accelerate access further (even before consulting hash table)





# Summary

---

- Advantages of Paging – shared memory, avoid external fragmentation
- Requirements for shared paging – valid / invalid bit
  
- Where to store the Page Table? In the memory
  - Cost – multiple memory accesses for each data access
  - Optimization for faster access – use TLB (associative cache)
  
- Structure of the Page Table in memory
  - Page Table can be too large for contiguous memory allocation
  - Hierarchical multi-level Page Table
    - ▶ Store pages of page table (POPT)
    - ▶ Does not scale well as logical address space becomes larger
  - Hashed Page Table
  - Inverted Page Table

